Metrics across the lifecycle: a framework for evaluating NIH biomedical data ecosystems

This manuscript (permalink) was automatically generated from seandavi/2026-cfde-metrics-manuscript@850acfb on March 17, 2026.

Authors

- **Ethan M. Lange**
  0000-0001-7075-4287
  Department of Biomedical Informatics, University of Colorado Anschutz Medical Campus · Funded by Grant U54OD036472

- **Vikram Adithya Ganesh**
  0009-0009-2056-2696 · techvik
  Department of Biomedical Informatics, University of Colorado Anschutz Medical Campus · Funded by Grant U54OD036472

- **Faisal Alquaddoomi**
  XXXX-XXXX-XXXX-XXXX · falquaddoomi
  Department of Biomedical Informatics, University of Colorado Anschutz Medical Campus · Funded by Grant U54OD036472

- **David Mayer**
  XXXX-XXXX-XXXX-XXXX · the-mayer
  Department of Biomedical Informatics, University of Colorado Anschutz Medical Campus · Funded by Grant U54OD036472

- **Adriana Ivich**
  XXXX-XXXX-XXXX-XXXX · ivichadriana
  Department of Biomedical Informatics, University of Colorado Anschutz Medical Campus · Funded by Grant U54OD036472

- **Vince Rubinetti**
  XXXX-XXXX-XXXX-XXXX · vincerubinetti
  Department of Biomedical Informatics, University of Colorado Anschutz Medical Campus · Funded by Grant U54OD036472

- **Casey Greene** ✉
  0000-0001-8713-9213 · cgreene
  Department of Biomedical Informatics, University of Colorado Anschutz Medical Campus · Funded by Grant U54OD036472

- **Sean Davis** ✉
  0000-0002-8991-6458 · seandavis
  Department of Biomedical Informatics, University of Colorado Anschutz Medical Campus · Funded by Grant U54OD036472

✉ — Correspondence possible via GitHub Issues or email to Casey Greene <casey.s.greene@cuanschutz.edu>, Sean Davis <seandavi@gmail.com>.

Bigger Picture

The U.S. government spends hundreds of millions of dollars every year creating and maintaining large collections of biomedical research data — from genome sequences of hundreds of thousands of people to detailed maps of every cell type in the human body. The goal is to make these datasets openly available to scientists everywhere, accelerating discoveries that no single laboratory could achieve alone.

But how do we know whether this investment is paying off? At present, each funded program tracks its own success metrics — if it tracks them at all — making it nearly impossible to compare programs or to learn from what works. This is a significant missed opportunity. In the business world, organizations use structured “performance frameworks” to evaluate products, adjust strategies, and decide where to invest next. Publicly funded science deserves the same discipline.

In this paper, we describe three frameworks for evaluating biomedical data resources: one focused on how widely the data are used and scientifically cited, one on whether the data are well-documented and reusable by others, and one on whether the underlying infrastructure is reliable and financially sustainable. We argue that different metrics matter more at different stages of a project’s life — early on, data quality and infrastructure are paramount; later, scientific impact and long-term cost become critical. Standardizing a common evaluation language across all NIH data programs would enable better decisions about where to invest, how to sustain valuable resources, and when to gracefully retire those that have served their purpose.

Summary

The U.S. National Institutes of Health has invested billions of dollars in large-scale biomedical data resources, yet evaluation of these investments remains fragmented — each program tracks different metrics in different ways, preventing portfolio-level comparison and informed decision-making. We propose a structured approach built on three complementary prioritization frameworks: *public value* (user engagement, citations, funded grants), *scientific data quality* (metadata completeness, FAIR compliance, data dictionaries), and *operations and finance* (infrastructure reliability, cost sustainability). Critically, we argue that the relative importance of each framework shifts predictably across a data resource's lifecycle — from early infrastructure-building, through active use and impact, to long-term sustainability. Drawing on data from the NIH Common Fund Data Ecosystem (CFDE), we identify gaps in current measurement practice and propose concrete metrics and tools — including FAIR maturity models, GitHub analytics, and cost-forecasting frameworks — to close them.

Main text

Over the past decade, the U.S. National Institutes of Health (NIH) has made extraordinary investments in large-scale publicly available biomedical data resources. These include the Cancer Research Data Commons [1] (CRDC, NCI), the Cancer Genome Atlas (TCGA) [2], the Trans-Omics for Precision Medicine program [3] (TOPMed, NHLBI), the All-of-Us Research Program [4]. and . NIH has complemented these with cloud computing environments — the Cancer Genomics Cloud [5], BioData Catalyst [6], and AnVIL [7] — designed to make analysis at scale accessible to any researcher. The collective ambition is to democratize access to reusable data and accelerate scientific discovery.

The Common Fund Data Ecosystem [8] (CFDE) takes a slightly different approach. Building upon the data coordinating centers spanning 19 projects including GTEx [9], the Human Microbiome Project [10,11], and HuBMAP [12], the CFDE is a cross-cutting initiative that harmonizes metadata across these diverse datasets and provides a unified portal for discovery and access. The CFDE's mission is to create a scalable, sustainable infrastructure that enables researchers to find and reuse data across the Common Fund portfolio, with the ultimate goal of accelerating scientific discovery and improving human health. The ecosystem is organized into five specialized centers: the Data Resource Center (DRC), the Knowledge Center (KC), the Training Center (TC), the Cloud Workspace Implementation Center (CWIC), and the Integration and Coordination Center (ICC). The ICC bears primary responsibility for leading annual program-wide metric collection and facilitating evaluation efforts in partnership with NIH program staff, ensuring a continuous improvement cycle and building the evidence base for future science investments. With five centers and 19 projects, the CFDE is a microcosm of the broader NIH data ecosystem, making it an ideal case study for developing and testing federated evaluation frameworks.

Evaluation across the CFDE has historically been somewhat fragmented: each program or projects tracks slightly different metrics in different ways, making portfolio-level evaluation challenging. The business world has long solved analogous problems with structured prioritization frameworks that combine multiple performance metrics into coherent decision tools. Public science needs equivalent rigor. We contend that the "customer" is the research community, the "product" is reusable data and tools, and the measure of success is scientific discovery and improved human health. We describe three complementary prioritization frameworks–*public value*, *scientific data quality*, *operations and finance*–and suggest that their importance and prioritization should vary over a project lifecycle.

Performance metrics have long been used in the business world to optimize decision-making, with established frameworks ranging from quantitative scoring functions such as RICE (Reach, Impact, Confidence, Effort) [13/] to qualitative prioritization methods such as MoSCoW ("must-have", "should-have", "could-have", "won't have") [14]. Crucially, effective prioritization frameworks rely on metrics that are informative, feasible to collect, and economical to measure [15,16]. In publicly-funded research, the same principles apply: good metrics capture the use and impact of data resources, identify barriers to their success, and inform decisions about where to invest limited resources while being feasible to collect and not overly burdensome for project owners. Furthermore, we show how centralization and automation can make the collection of these metrics more efficient, reducing the burden on individual projects and enabling real-time monitoring of portfolio performance.

Three Frameworks, One Lifecycle

No single metric — and no single framework — can capture the full picture of a data resource's value or trajectory. We organize our approach around three complementary prioritization frameworks: *public value*, *scientific data quality*, and *operations and finance*. These frameworks overlap and sometimes conflict (more scientifically impactful datasets are often more expensive to create and

maintain), and weighted composite scoring (*Multi-Criteria Decision Analysis*) can balance them when needed [17,18]. Our central argument is that these frameworks are not equally relevant at all times: metrics are a function of a resource's lifecycle phase. In the **early term** — infrastructure-building and data collection — *scientific data quality* and *operations and finance* dominate. In the **mid term** — active use within a funded period — all three matter, with *public value* taking on increasing importance. In the **late, post-funding term** — sustainability and continued scientific relevance — *operations and finance* and *public value* become the primary lenses. Embedding this lifecycle logic into evaluation design, rather than applying uniform metrics across all phases, is the key practical contribution of this paper.

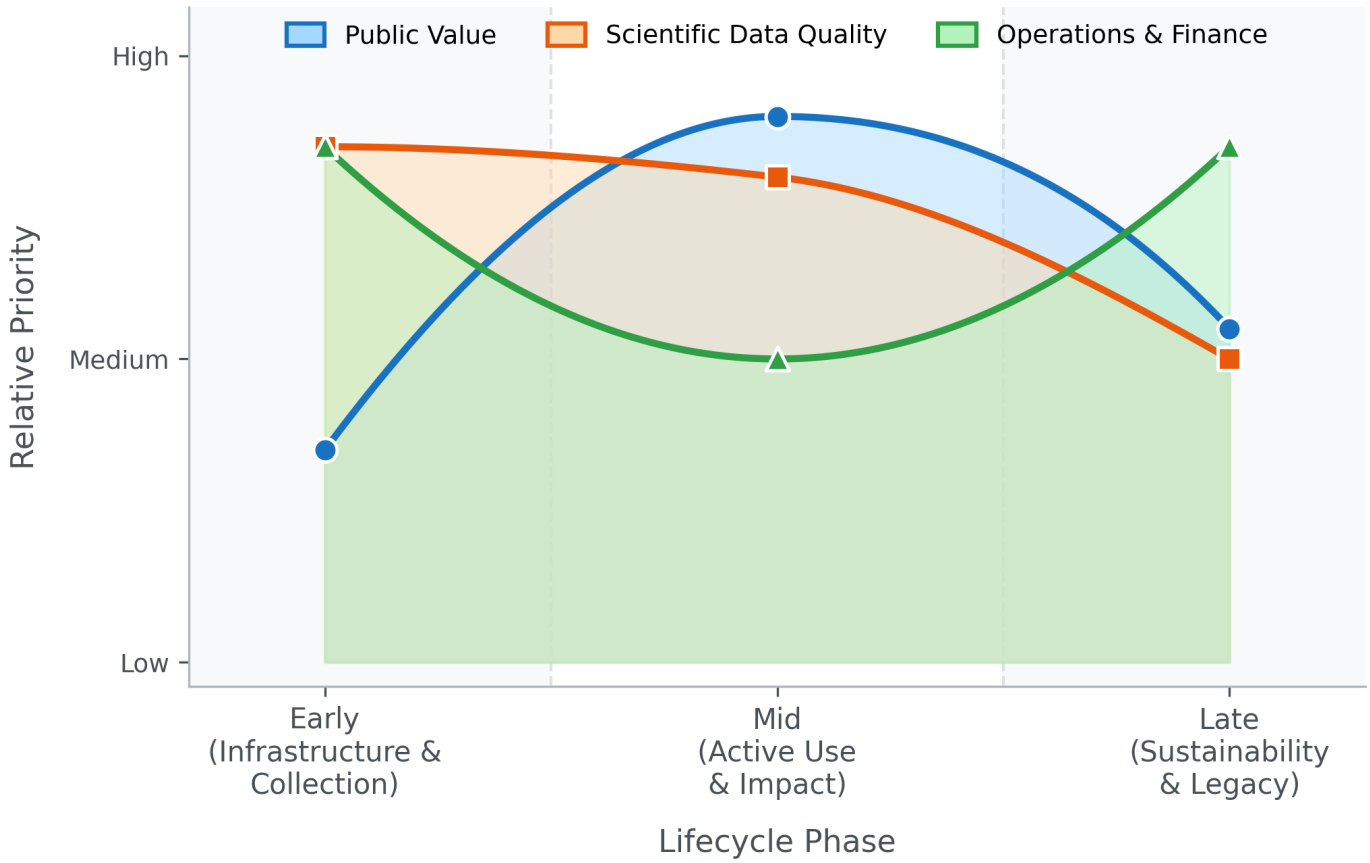

Figure 1: Schematic relative priority of the three metric and evaluation frameworks–public value, scientific data quality, and operations and finance–across the project lifecycle. Note that lower priority does **not** imply unimportant.

The *public value* framework captures the outputs and impacts of a data resource: user engagement, data downloads, funded grants, publications, and downstream scientific influence. The *scientific data quality* framework measures reusability and long-term value, guided by the FAIR principles (Findable, Accessible, Interoperable, Reusable) [19,20] and encompassing the metadata, documentation, and standardization needed for data to be reliably reused. The *operations and finance* framework covers the infrastructure and cost considerations — from server uptime and compute performance to the sustainability of funding beyond the initial grant period. For each framework, we describe key metrics, their strengths and limitations, and highlight what current measurement approaches are missing.

Public value

The *public value* framework encompasses user behavior and scientific impact. Website engagement metrics — *page views*, *time on page*, *actions/triggers*, *priority link clicks*, and *new vs. returning users* — are easy to collect and track trends over time, but they provide limited insight into actual data use. Combining *time on page* with *actions/triggers* offers a more informative measure of engagement. A key limitation is that bots can skew counts, and VPN use in academic settings confounds new-versus-

returning user tracking. These metrics are typically collected via platforms such as Google Analytics [21/] or Matomo [22/] (direct measurement) or estimated via Semrush [23/] or Similarweb [24/]; note that some European jurisdictions restrict use of US-based analytics platforms.

More direct evidence of data engagement comes from *downloads* and *compute jobs*. Downloads are a simple indicator of interest but cannot distinguish actual usage from passive acquisition, nor track data sharing among multiple users of the same file. Cloud-hosted datasets address these limitations: compute jobs directly measure analysis activity, reduce sharing ambiguity, and can be monitored in near real-time (*trend downloads*, *trend compute jobs*). Projects that include tool development — such as GlyGen [25] or Bridge2AI [26] — also track *tool usage* (online) and *tool downloads*, with online usage being more accurately quantified.

User behavior metrics do not capture scientific impact. The strongest proxies are publications and citations, accessible via PubMed [27/], Scopus [28/], Google Scholar [29/], and Web of Science [30/]. *Number citations* is a direct measure of data influence; the initial GTEx paper [9], for example, had accumulated 8,630 citations as of January 2026. Citation impact can be weighted by the influence of each citing work — either by that paper's own citation count (*number secondary citations*) or by journal-level metrics (*citing impact score*). A composite score weighted by secondary citations (*citations + secondary citations*) captures cumulative downstream impact that simple citation counts miss, and citation networks can also reveal new cross-disciplinary collaborations enabled by data sharing [**doi:10.1038/s42254-021-00368-x?**]. Key limitations of citation-based metrics are that they lack context (a citation does not convey how central the dataset was to a study), and they lag by months to years relative to actual use. Tracking the trend in citations over time (*trend citations*) helps identify whether a resource's scientific influence is growing, plateauing, or declining.

A fundamental challenge for citation-based evaluation is that data citation is not yet a universal standard in scientific publishing. Studies estimate that fewer than 30% of articles performing secondary analysis provide a formal reference to the dataset used [31]. This "citation gap" means that standard bibliometric approaches systematically underestimate data reuse. Evaluators must therefore supplement citation tracking with multi-pronged strategies: searching for unique accession numbers (e.g., GEO accession numbers, NCBI BioProject IDs) in the full text of publications, mining Methods sections for mentions of specific repositories, and conducting case-study analyses of high-value datasets. Bibliometric analysis of NeuroMorpho.Org, for example, demonstrated that sharing neural morphology data leads to a significant increase in citations for the original authors, with a reuse rate that remains constant for at least 16 years — effectively doubling peer-reviewed discoveries in the field [31].

Beyond usage and citation metrics, user-centered measures of *perceived value* offer complementary insight into how well a data platform serves its community. Research integrating the Information System Success Model, the Technology Acceptance Model, and Consumer Perceived Value Theory identifies four entropy-weighted metrics that explain the majority of perceived value: *number of datasets* (breadth of coverage), *data timeliness* (currency for emergent challenges), *search comprehensiveness* (ability to return all relevant records across federated sources), and *system responsiveness* (loading times and API stability) [32]. These metrics can be assessed through structured surveys and complement the behavioral measures captured by web analytics.

Scientific data quality

The most profound long-term measure of public value is *practice change* — a fundamental shift in how science is conducted. Concrete indicators include the adoption rate of cloud-based workspaces versus local data downloads (a proxy for computational paradigm shift), cross-disciplinary co-authorship networks across Common Fund programs (a measure of new collaborative partnerships), and the geographic and institutional diversity of the user base, particularly the participation of early-

stage investigators from under-resourced institutions. These indicators capture whether a data ecosystem is achieving its ultimate goal of democratizing access and accelerating discovery.

The *scientific data quality* framework addresses whether a resource's data can actually be found, understood, and reused by others. Poor data quality and poor annotation are leading barriers to data reuse. At the broadest level, metadata metrics document the study — its name, description, dates, funding, contributors, and website — and should be uniform across projects to enable portfolio-level comparisons. At a finer grain, data-level metrics capture storage format, dataset size, data types (clinical, genomic, metabolomic, etc.), measurement platforms (e.g., Olink Explore HT for proteomics), biospecimen handling, and collection conditions (e.g., fasting status). At the finest grain, dataset-level metrics include sample sizes, variable counts, demographic summaries, and QC statistics.

For evaluation purposes, summary metrics spanning all three levels are more tractable than cataloging every study-specific variable. *Metadata completeness* counts the presence of required fields and can be assessed programmatically. *Standardized metadata* assesses whether fields conform to community guidelines — a harder problem that often requires human review. *Quality-control (QC) compliance* and *data uniqueness* are similarly difficult to automate and frequently require trained evaluators. Two additional high-value metrics are the presence (*available data dictionary*) and quality (*quality data dictionary*) of a data dictionary that documents variable definitions, units, timing, and transformations; the former can be detected by script, the latter requires domain expertise. The Standard Evaluation Framework for Large Health Care Data identifies 49 specific attributes grouped into three domains — the collection process, the data itself, and data use — that provide a comprehensive approach to assessing fitness for purpose [33]. Two attributes from this framework are particularly relevant for federated ecosystems like the CFDE: *representativeness*, the extent to which a dataset generalizes beyond its source population (critical for avoiding false conclusions), and *linkability*, the extent to which data can be integrated with external sources — a paramount concern for an ecosystem whose mission is to enable queries across diverse programs.

A structured approach to assessing these qualities is provided by FAIR maturity models. The FAIRplus Dataset Maturity (FAIRplus-DSM) model , developed for the life sciences, defines six levels ranging from Level 0 (no unique identifiers, unshared local storage) through Level 5 (provenance encoded in cross-community languages, enterprise-grade lifecycle management). Levels 1–2 cover persistent identifiers and indexed metadata; Levels 3–4 require adherence to community reporting standards (e.g., MIAME) and semantic representation using linked data. The Research Data Alliance (RDA) further classifies FAIR indicators by priority — Essential (globally unique persistent identifiers, open access protocols), Important (machine-understandable knowledge representation), and Useful (automated data access by software agents) — providing a tiered roadmap for incremental improvement. Automated tools such as the FAIR Evaluator can test digital resources against these indicators at scale, making continuous *FAIR compliance* monitoring feasible for large portfolios.

Operations and finance

The *operations and finance* framework addresses whether the infrastructure supporting a data resource is reliable, performant, and sustainable. On the operations side, server-level metrics — *uptime*, *page load time*, *latency*, *server errors*, and *client errors* — can be captured via analytics platforms and log-file scripts. Compute-level metrics (*download speed*, *cpu/gpu usage*, *memory usage*, *client queue times*, *job run times*) identify performance bottlenecks. These metrics carry inherent variability due to time-of-day load, job size, and client-side network conditions, so monitoring trends rather than point-in-time values is most informative.

For projects built on open-source software — as many CFDE components are — GitHub repository analytics provide an additional operational window. Development velocity (commit frequency, pull-

request review latency, merge frequency) reflects team capacity and responsiveness to community contributions. Codebase health can be assessed by tracking stale pull requests and unresolved issues, which can signal declining institutional support or funding gaps. Contributor dynamics — distinguishing core team commits from external community contributions — indicate whether training and outreach efforts are successfully broadening adoption. For a more structured assessment of software components, the ELIXIR "Top 10 Metrics for Life Science Software Good Practices" provides a prioritized framework ranked by importance (impact on sustainability) and implementability (ease of generation) [34]. These metrics — spanning version control, software discoverability, automated builds, test data availability, and documentation — offer a reproducible scorecard for evaluating the software that underpins data infrastructure, and can be integrated into evaluation dashboards via automated collection frameworks.

On the finance side, *funding cost* is the most accessible metric, but a full cost picture requires decomposing it into components: *sample collection*, *sample measurement*, *QC*, *data storage*, *data preservation* (beyond the grant period), *computing hardware*, *server maintenance*, *cloud-based computing*, and *personnel* costs. The National Academies framework for forecasting biomedical information resource costs identifies six cost-driver categories — data content, metadata, labor, infrastructure, compliance, and governance — and notes that labor (curation and user support) consistently represents the largest single cost element. A critical but often overlooked question is the break-even point for data sharing: mathematical models suggest that the efficiency gains to the research community from reusing shared data do not always outweigh the curation costs borne by generators, highlighting the need for selective investment in high-value datasets. More broadly, evaluating the return on investment (ROI) of a data ecosystem requires distinguishing between direct benefits (cost savings from economies of scale in data generation and curation) and extended benefits (scientific discoveries and improved clinical outcomes enabled by data reuse). Cloud computing introduces additional complexity, as costs are provider-specific, usage-dependent, and subject to change unpredictably. A summary of all metrics across frameworks is provided in Table 1.

Table 1. Quick summary of included metrics and

Metric	Difficulty	How Collected	Value	Lifecycle Phase	Notes
User Behavior Framework					
page views	Easy	Google Analytics	Low	Mid	A minimum indicator of interest.
time on page	Easy	Google Analytics	Low	Mid	A minimum indicator of interest.
actions/triggers	Easy	Google Analytics (Tag Manager)	Medium	Mid	A moderate indicator of interest.
priority link clicks	Easy	Google Analytics (Tag Manager)	Medium	Mid	A moderate indicator of interest.
new vs returning users	Medium	Google Analytics	Medium	Mid	A moderate indicator of continuing interest.

Metric	Difficulty	How Collected	Value	Lifecycle Phase	Notes
downloads	Easy	Google Analytics or Server-side logs	Medium-High	Mid	A strong indicator of interest. Downloading not equivalent to usage.
compute jobs	Easy	Compute scheduler logs or cloud environment logs	High	Mid	A direct indicator of interest.
tool downloads	Easy	Google Analytics or server-side logs	Medium-High	Mid	A strong indicator of interest.
tool usage	Easy	Google Analytics or API logs	High	Mid	A direct indicator of interest.
trend downloads	Medium	Google Analytics or server-side logs BI tools for visualization	Medium-High	Mid–Late	A strong indicator of change in interest over time.
trend compute jobs	Medium	Compute scheduler logs or cloud environment logs BI tools for visualization	High	Mid–Late	A direct indicator of change in interest over time.
number grants	Hard	NIH RePORTER Self-reporting	High	Mid–Late	Relevant grant data not publicly available.
number presentations	Hard	Self-reporting	Medium	Mid–Late	Data difficult to track and often not available in sufficient detail.
number citations	Easy	Google Scholar or CrossRef or other similar platforms	High	Mid–Late	A direct indicator of data value. Can miss downstream cumulative value.
number secondary citations	Medium		High	Late	An indirect measure of “downstream” value of initial publications.
citations + secondary citations	Medium		High	Late	A measure of cumulative impact over time of study data.
trend citations	Medium		High	Late	Measures change in impact of data over time. Subject to time lag in publications.

Metric	Difficulty	How Collected	Value	Lifecycle Phase	Notes
data citation rate	Hard	Full-text search for accession numbers, Methods section mining	High	Mid–Late	<30% of secondary analyses formally cite datasets. Multi-pronged tracking required.
perceived value (user surveys)	Medium	Structured surveys (ISSM/TAM framework)	High	Mid–Late	Entropy-weighted: dataset breadth, timeliness, search comprehensiveness, responsiveness.
cloud adoption rate	Medium	Cloud workspace logs vs. download logs	Medium-High	Mid–Late	Proxy for computational paradigm shift (cloud vs. local analysis).
cross-disciplinary collaboration	Hard	Co-authorship network analysis	High	Late	Measures new partnerships fostered across Common Fund programs.
user diversity (geographic/institutional)	Medium	User registration data, IAM logs	Medium-High	Mid–Late	Measures democratization of access, especially for under-resourced institutions.
Scientific Data Quality Framework					
metadata completeness	Medium	Metadata validator, Schema validator	Medium-High	Early	Critical for drawing initial interest in study.
standardized metadata	Hard	Requires human-in-the-loop evaluation.	High	Early	Study-specific requirements.
quality-control (QC) compliance	Hard	Requires human-in-the-loop evaluation.	High	Early	Study-specific requirements.
data uniqueness	Hard	Requires human-in-the-loop evaluation.	Medium	Early	Subjective measure.
FAIR compliance	Hard	Requires human-in-the-loop evaluation.	High	All	Subjective measure.
available data dictionary	Medium	Checks on repository or documentation	Medium	Early	An essential component for a study to have. Does not guarantee quality.

Metric	Difficulty	How Collected	Value	Lifecycle Phase	Notes
quality data dictionary	Hard	Requires human-in-the-loop evaluation.	High	Early	A metric of quality of the data dictionary.
representativeness	Hard	Demographic and population metadata analysis	High	Early	Whether dataset generalizes beyond source population. Critical for avoiding false conclusions.
linkability	Hard	Schema and identifier compatibility assessment	High	All	Extent to which data integrates with external sources. Paramount for federated ecosystems.
***Operations and Finance* Framework**					
uptime	Easy	Uptime monitoring tool based on infrastructure/cloud used	Medium-Low	Mid–Late	Measure of server stability.
page load time	Easy	Google Analytics	Medium	Mid–Late	Measure of server performance.
latency	Easy	Google Analytics (not real-time)	Medium	Mid–Late	Measure of server performance.
server errors	Easy	Server-side logs or Application Performance Monitoring (APM) tools	Medium	Mid–Late	Measure of server stability and performance.
client errors	Medium	Application logs or Application Performance Monitoring (APM) tools	Low	Mid–Late	Measure of errors on client end. Possible not addressable.
download speed	Easy	Server-side logs or Application Performance Monitoring (APM) tools	Low-Medium	Mid–Late	Download times can be function of client connection speeds.
cpu/gpu usage	Easy	Cloud/infrastructure logs or dashboards	High	Mid–Late	Important to evaluate performance and additional needs.

Metric	Difficulty	How Collected	Value	Lifecycle Phase	Notes
memory usage	Easy	Cloud/infrastructure logs or dashboards	High	Mid–Late	Important to evaluate performance and additional needs.
number of users	Easy	Google Analytics or IAM logs for granular details (if available)	Medium	Mid–Late	Indirect measure of computational load/burden.
client queue times	Easy	Cloud/infrastructure logs or dashboards	Medium-High	Mid–Late	Long queue times can result in less usage.
job run times	Easy	Cloud/infrastructure compute logs or dashboards	Medium	Mid–Late	Excessive long job run times can result in less usage.
software good practices score	Medium	ELIXIR Top 10 framework; automated collection via PIsCO	High	All	Covers version control, discoverability, automated builds, test data, documentation.
funding cost	Easy		High	All	Total cost of study. Critical element in evaluating study feasibility and sustainability.
sample collection costs	Medium		Medium	Early	Cost typically associated with early study period. Cost does not impact sustainability.
sample measurement costs	Medium		Medium	Early	Cost typically associated with early study period. Cost does not impact sustainability.
QC costs	Medium		Medium	Early	Cost typically associated with early study period. Cost does not impact sustainability.
data storage costs	Hard		High	All	Costs and requirements can fluctuate over study.

Metric	Difficulty	How Collected	Value	Lifecycle Phase	Notes
data preservation costs	Hard		High	Late	Costs and requirements can fluctuate over study. Difficult to assess years in advance.
computing hardware costs	Medium		Medium	Early	Costs often associated with initial study period, but replacement costs are possible.
server maintenance costs	Hard		High	Mid–Late	Costs can fluctuate. Major source of cost in sustainability period.
cloud-based computing costs	Hard	Cloud Cost Explorer tools or	High	Mid–Late	Costs are highly variable and subject to constant change (for both host and clients). Especially difficult to assess years in advance.
personnel costs	Medium		High	All	Can be measured in money or time. Requires committed time with relevant expertise.

Synthesis and Recommendations

The three frameworks described here — *public value*, *scientific data quality*, and *operations & finance* — are complementary rather than competing, and their relative importance shifts predictably across the data resource lifecycle. Early-phase investments should prioritize the scientific data quality and operations frameworks: getting metadata right, establishing FAIR-compliant structures, and standing up reliable infrastructure are prerequisites for everything that follows. Mid-phase evaluation should expand to all three frameworks, with particular attention to public value metrics as user communities engage with the data. Late-phase and sustainability-focused evaluation should center on the public value and operations & finance frameworks, asking whether the resource continues to generate scientific impact and whether its costs can be justified in the absence of original funding.

No framework operates in isolation, and quantitative metrics alone cannot capture the full picture. Mixed-method approaches — combining usage statistics and citation data with structured user interviews, focus groups, and surveys — are essential for understanding *why* certain resources succeed or fail, not just *what* is happening [35]. Two established designs are particularly suited to ecosystem evaluation: *sequential explanatory* designs, in which quantitative data (e.g., utilization rates) are collected first and then followed by qualitative interviews or focus groups to explain observed patterns, and *convergent* designs, in which both data types are collected simultaneously and compared at the interpretation stage to triangulate findings. Logic models that explicitly connect

inputs (funding, data, personnel) to activities (harmonization, training, curation), outputs (portals, trained users, tools), and outcomes (publications, funded grants, practice change) [36] provide the evaluative scaffolding needed to identify which elements of an ecosystem are working. These models should be treated as living documents, revised as programs transition from pilot to full implementation.

A critical long-term challenge is closing the "know-do gap" — the distance between the information provided by metrics and actual decision-making. Evidence from healthcare quality improvement suggests that the act of repeated, structured measurement itself drives behavioral change among data producers and managers, encouraging improved metadata practices and FAIR compliance [37]. NIH's 2023 Data Management and Sharing Policy creates a structural opportunity to establish pre-policy baselines for sharing and reuse, enabling rigorous before-and-after assessment of whether mandates translate into measurable ecosystem improvements.

Implementation of these evaluation frameworks can be accelerated through pilot "Driver Projects" — an approach pioneered by the Global Alliance for Genomics and Health (GA4GH) [38/], in which selected projects test new standards and metrics in real-world settings and provide deep qualitative feedback on their utility. For the CFDE, identifying Driver Projects across centers would generate practical evidence about which metrics are most informative and feasible before committing to portfolio-wide deployment. Ultimately, the long-term measure of evaluation success is whether the ecosystem achieves practice change: a measurable shift toward data-driven discovery, broader cross-disciplinary collaboration, and more equitable access to computational resources.

Looking forward, the field needs shared dashboards that integrate technical health indicators (FAIR scores, GitHub analytics, software good practices scores, uptime), impact metrics (citation trends, compute job volumes, perceived value surveys), and financial sustainability projections into a coherent, real-time view of portfolio performance. Standardizing a core set of metrics across all NIH Common Fund programs — while allowing program-specific extensions — would enable the portfolio-level comparisons that are currently impossible. The frameworks presented here, grounded in real evaluation data from the CFDE, offer a starting point for that standardization.

References

1. **A Comprehensive Infrastructure for Big Data in Cancer Research: Accelerating Cancer Research and Precision Medicine**
Izumi V Hinkson, Tanja M Davidsen, Juli D Klemm, Ishwar Chandramouliswaran, Anthony R Kerlavage, Warren A Kibbe
Frontiers in Cell and Developmental Biology (2017-09-21) https://doi.org/gpf2r8
DOI: 10.3389/fcell.2017.00083 · PMID: 28983483 · PMCID: PMC5613113

2. **The Cancer Genome Atlas Pan-Cancer analysis project.**
, John N Weinstein, Eric A Collisson, Gordon B Mills, Kenna RMills Shaw, Brad A Ozenberger, Kyle Ellrott, Ilya Shmulevich, Chris Sander, Joshua M Stuart
Nature genetics (2013-10) https://www.ncbi.nlm.nih.gov/pubmed/24071849
DOI: 10.1038/ng.2764 · PMID: 24071849 · PMCID: PMC3919969

3. **Sequencing of 53,831 diverse genomes from the NHLBI TOPMed Program**
Daniel Taliun, Daniel N Harris, Michael D Kessler, Jedidiah Carlson, Zachary A Szpiech, Raul Torres, Sarah AGagliano Taliun, André Corvelo, Stephanie M Gogarten, Hyun Min Kang, … Gonçalo R Abecasis
Nature (2021-02-10) https://doi.org/gh4bkp
DOI: 10.1038/s41586-021-03205-y · PMID: 33568819 · PMCID: PMC7875770

4. **The "All of Us" Research Program** *New England Journal of Medicine* (2019-08-15) https://doi.org/ggjdzj
DOI: 10.1056/nejmsr1809937 · PMID: 31412182

5. **The Cancer Genomics Cloud: Collaborative, Reproducible, and Democratized—A New Paradigm in Large-Scale Computational Research**
Jessica W Lau, Erik Lehnert, Anurag Sethi, Raunaq Malhotra, Gaurav Kaushik, Zeynep Onder, Nick Groves-Kirkby, Aleksandar Mihajlovic, Jack DiGiovanna, Mladen Srdic, … Brandi Davis-Dusenbery
Cancer Research (2017-10-31) https://doi.org/gcjxm2
DOI: 10.1158/0008-5472.can-17-0387 · PMID: 29092927

6. **The NHLBI BioData Catalyst**
Lung National Heart
Zenodo (2020) https://doi.org/hbswtd
DOI: 10.5281/zenodo.3822858

7. **Inverting the model of genomics data sharing with the NHGRI Genomic Data Science Analysis, Visualization, and Informatics Lab-space**
Michael C Schatz, Anthony A Philippakis, Enis Afgan, Eric Banks, Vincent J Carey, Robert J Carroll, Alessandro Culotti, Kyle Ellrott, Jeremy Goecks, Robert L Grossman, … Jason Walker
Cell Genomics (2022-01) https://doi.org/gn7rf9
DOI: 10.1016/j.xgen.2021.100085 · PMID: 35199087 · PMCID: PMC8863334

8. **Making Common Fund data more findable: catalyzing a data ecosystem**
Amanda L Charbonneau, Arthur Brady, Karl Czajkowski, Jain Aluvathingal, Saranya Canchi, Robert Carter, Kyle Chard, Daniel JB Clarke, Jonathan Crabtree, Heather H Creasy, … Owen White
GigaScience (2022) https://doi.org/gsxq6b
DOI: 10.1093/gigascience/giac105 · PMID: 36409836 · PMCID: PMC9677336

9. **The Genotype-Tissue Expression (GTEx) project**

John Lonsdale, Jeffrey Thomas, Mike Salvatore, Rebecca Phillips, Edmund Lo, Saboor Shad, Richard Hasz, Gary Walters, Fernando Garcia, Nancy Young, … Helen F Moore
Nature Genetics (2013-05-29) https://doi.org/gd5z68
DOI: 10.1038/ng.2653 · PMID: 23715323 · PMCID: PMC4010069

10. **Structure, function and diversity of the healthy human microbiome** *Nature* (2012-06) https://doi.org/hzm
DOI: 10.1038/nature11234 · PMID: 22699609 · PMCID: PMC3564958

11. **A framework for human microbiome research** *Nature* (2012-06) https://doi.org/f32svq
DOI: 10.1038/nature11209 · PMID: 22699610 · PMCID: PMC3377744

12. **The human body at cellular resolution: the NIH Human Biomolecular Atlas Program.**
Nature (2019-10-09) https://www.ncbi.nlm.nih.gov/pubmed/31597973
DOI: 10.1038/s41586-019-1629-x · PMID: 31597973 · PMCID: PMC6800388

13. **RICE: Simple prioritization for product managers**
Sean McBride
The Intercom Blog (2018-01-05) https://www.intercom.com/blog/rice-simple-prioritization-for-product-managers/

14. **CASE method fast track: a RAD approach**
Dai Clegg, Richard Barker
Addison-Wesley Pub. Co (1994)
ISBN: 9780201624328

15. **The balanced scorecard: translating strategy into action**
Robert S Kaplan, David P Norton
Harvard Business School Press (1996)
ISBN: 9780875846514

16. **The performance prism: the scorecard for measuring and managing business success**
Andy D Neely, Chris Adams, Mike Kennerley
Financial Times Prentice Hall (2002)
ISBN: 9780273653349

17. **Decisions with multiple objectives: preferences and value tradeoffs**
Ralph L Keeney, Howard Raiffa
Wiley (1976)
ISBN: 9780471465102

18. **Advances in information security management & small systems security: IFIP TC11 WG11.1/WG11.2 Eighth Annual Working Conference on Information Security Management & Small Systems Security, September 27-28, 2001, Las Vegas, Nevada, USA**
Jan HP Eloff (editor)
Kluwer Academic Publishers (2001)
ISBN: 9780792375067

19. **The FAIR Guiding Principles for scientific data management and stewardship**
Mark D Wilkinson, Michel Dumontier, IJsbrand Jan Aalbersberg, Gabrielle Appleton, Myles Axton, Arie Baak, Niklas Blomberg, Jan-Willem Boiten, Luiz Bonino da Silva Santos, Philip E Bourne, … Barend Mons
Scientific Data (2016-03-15) https://doi.org/bdd4
DOI: 10.1038/sdata.2016.18 · PMID: 26978244 · PMCID: PMC4792175

20. **A design framework and exemplar metrics for FAIRness**

Mark D Wilkinson, Susanna-Assunta Sansone, Erik Schultes, Peter Doorn, Luiz Olavo Bonino da Silva Santos, Michel Dumontier
Scientific Data (2018-06-26) https://doi.org/gfkrrt
DOI: 10.1038/sdata.2018.118 · PMID: 29944145 · PMCID: PMC6018520

21. **Analytics Tools & Solutions for Your Business - Google Analytics**
Google Marketing Platform
https://marketingplatform.google.com/about/analytics/

22. **Matomo - The Google Analytics alternative that protects your data**
Analytics Platform - Matomo
https://matomo.org/

23. **Semrush: Data-Driven Marketing Tools to Grow Your Business**
Semrush
https://www.semrush.com/

24. **Similarweb: AI-Powered Digital Data Intelligence Solutions**
Similarweb
https://www.similarweb.com/

25. **GlyGen: Computational and Informatics Resources for Glycoscience**
William S York, Raja Mazumder, Rene Ranzinger, Nathan Edwards, Robel Kahsay, Kiyoko F Aoki-Kinoshita, Matthew P Campbell, Richard D Cummings, Ten Feizi, Maria Martin, … Wenjin Zhang
Glycobiology (2019-10-16) https://doi.org/hbswtb
DOI: 10.1093/glycob/cwz080 · PMID: 31616925 · PMCID: PMC7335483

26. **Reproducible generation of human liver organoids (HLOs) on a pillar plate platform via microarray 3D bioprinting**
Sunil Shrestha, Vinod Kumar Reddy Lekkala, Prabha Acharya, Soo-Yeon Kang, Manav Goud Vanga, Moo-Yeal Lee
openRxiv (2024-03-13) https://doi.org/hbswtc
DOI: 10.1101/2024.03.11.584478 · PMID: 38559126 · PMCID: PMC10979895

27. **PubMed**
PubMed
https://pubmed.ncbi.nlm.nih.gov/

28. https://www.scopus.com

29. **Google Scholar** https://scholar.google.com/

30. **Clarivate** https://access.clarivate.com/login?app=wos&alternative=true&shibShireURL=https%3A%2F%2Fwww.webofknowledge.com%2F%3Fauth%3DShibboleth&shibReturnURL=https%3A%2F%2Fwww.webofknowledge.com%2F&roaming=true

31. **Bibliometric analysis of neuroscience publications quantifies the impact of data sharing**
Herve Emissah, Bengt Ljungquist, Giorgio A Ascoli
openRxiv (2023-09-13) https://doi.org/hbsw6q
DOI: 10.1101/2023.09.12.557386 · PMID: 37745378 · PMCID: PMC10515804

32. **Medical Science Data Value Evaluation Model: Mixed Methods Study**
Dandan Wang, Yaning Liu
JMIR Medical Informatics (2025-08-21) https://doi.org/hbsw6s
DOI: 10.2196/63544 · PMID: 40839787 · PMCID: PMC12369987

33. **A Standard Framework for Evaluating Large Health Care Data and Related Resources**
Suad El Burai Felix, Hussain Yusuf, Matthew Ritchey, Sebastian Romano, Gonza Namulanda, Natalie Wilkins, Tegan K Boehmer
MMWR Supplements (2024-05-09) https://doi.org/hbsw6r
DOI: 10.15585/mmwr.su7303a1 · PMID: 38713639 · PMCID: PMC11078514

34. **Top 10 metrics for life science software good practices**
Haydee Artaza, Neil Chue Hong, Manuel Corpas, Angel Corpuz, Rob WW Hooft, Rafael C Jiménez, Brane Leskošek, Brett G Olivier, Jan Stourac, Radka Svobodová Vařeková, ... Daniel Vaughan
F1000Research (2016-08-16) https://doi.org/gfvpq7
DOI: 10.12688/f1000research.9206.1 · PMID: 27635232 · PMCID: PMC5007752

35. **Purposeful Sampling for Qualitative Data Collection and Analysis in Mixed Method Implementation Research**
Lawrence A Palinkas, Sarah M Horwitz, Carla A Green, Jennifer P Wisdom, Naihua Duan, Kimberly Hoagwood
Administration and Policy in Mental Health and Mental Health Services Research (2013-11-06) https://doi.org/f7nxzd
DOI: 10.1007/s10488-013-0528-y · PMID: 24193818 · PMCID: PMC4012002

36. https://www.wkkf.org/resource-directory/resources/2004/01/logic-model-development-guide

37. **NOT-OD-21-013: Final NIH Policy for Data Management and Sharing**
https://grants.nih.gov/grants/guide/notice-files/NOT-OD-21-013.html

38. **Strategic Road Map** https://www.ga4gh.org/about-us/strategic-road-map/